

```
In [2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

```
In [3]: df=pd.read_csv(r"C:\Users\HP\Downloads\fiat500_VehicleSelection_Dataset (2).csv")
df
```

```
Out[3]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



```
In [4]: df.head()
```

```
Out[4]:
```


	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700



```
In [5]: df.tail()
```

```
Out[5]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
1533	1534	sport	51	3712	115280	1	45.069679	7.70492
1534	1535	lounge	74	3835	112000	1	45.845692	8.66687
1535	1536	pop	51	2223	60457	1	45.481541	9.41348
1536	1537	lounge	51	2557	80750	1	45.000702	7.68227
1537	1538	pop	51	1766	54276	1	40.323410	17.56827



```
In [6]: df.shape
```

```
Out[6]: (1538, 9)
```

```
In [7]: df.describe
```

```
Out[7]: <bound method NDFrame.describe of
```

	ID	model	engine_power	age_in_da
ys	km	previous_owners		
0	1	lounge	51	882
1	2	pop	51	1186
2	3	sport	74	4658
3	4	lounge	51	2739
4	5	pop	73	3074
...
1533	1534	sport	51	3712
1534	1535	lounge	74	3835
1535	1536	pop	51	2223
1536	1537	lounge	51	2557
1537	1538	pop	51	1766


	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

```
[1538 rows x 9 columns]>
```

```
In [8]: df.isna().any()
```

```
Out[8]: ID                False
model                  False
engine_power          False
age_in_days           False
km                    False
previous_owners        False
lat                   False
lon                   False
price                 False
dtype: bool
```

```
In [9]: sns.lmplot(x='km',y='price',data=df,order=2,ci=None)
plt.show()
```




```
In [10]: x=np.array(df['km']).reshape(-1,1)
y=np.array(df['price']).reshape(-1,1)
```

```
In [11]: df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
#splitting data into train and test
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


0.741062312488328

```
In [12]: y_pred=regr.predict(x_test)

plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [13]: df500=df[:][:500]
sns.lmplot(x="km",y="price",data=df500,order=1,ci=None)
plt.show()
```



```
In [14]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [15]: model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.741062312488328

```
In [16]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [19]: converter={"model":{"sport":1,"lounge":2,"pop":3}}
df=df.replace(converter)
df
```


```
Out[19]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	2	51	882	25000	1	44.907242	8.611560
1	2	3	51	1186	32500	1	45.666359	12.241890
2	3	1	74	4658	142228	1	45.503300	11.417840
3	4	2	51	2739	160000	1	40.633171	17.634609
4	5	3	73	3074	106880	1	41.903221	12.495650
...
1533	1534	1	51	3712	115280	1	45.069679	7.704920
1534	1535	2	74	3835	112000	1	45.845692	8.666870
1535	1536	3	51	2223	60457	1	45.481541	9.413480
1536	1537	2	51	2557	80750	1	45.000702	7.682270
1537	1538	3	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



```
In [20]: plt.figure(figsize = (10, 10))
sns.heatmap(df.corr(), annot = True)
plt.show()
```



```
In [21]: features=df.columns[0:1]
target=df.columns[-1]
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X_train is (1153, 1)

The dimension of X_test is (385, 1)

```
In [22]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.00310286926477088

The test score for lr model is -0.008405634316406507

```
In [23]: #Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.0031026398591535997


The test score for ridge model is -0.008307809466001403


```
In [24]: plt.figure(figsize=(10,10))
```

```
Out[24]: <Figure size 1000x1000 with 0 Axes>
```

```
<Figure size 1000x1000 with 0 Axes>
```

```
In [26]: plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=7,  
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,  
plt.xticks(rotation=90)  
plt.legend()  
plt.show()
```



```
In [27]: #Lasso regression model  
print("\nLasso Model: \n")  
lasso = Lasso(alpha = 10)  
lasso.fit(X_train,y_train)  
train_score_ls =lasso.score(X_train,y_train)  
test_score_ls =lasso.score(X_test,y_test)  
print("The train score for ls model is {}".format(train_score_ls))  
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.003075838461310987

The test score for ls model is -0.007367578602064606

```
In [37]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[37]: <Axes: >




```
In [29]: #Using the Linear CV model  
from sklearn.linear_model import LassoCV  
#Lasso Cross validation  
lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).  
#score  
print(lasso_cv.score(X_train, y_train))  
print(lasso_cv.score(X_test, y_test))
```

0.0031025989567363688

-0.008299466692577973

In [34]: *#plot size*

```
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6,color='red')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



```
In [38]: #Using the Linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.0031026398591535997

The train score for ridge model is -0.008307809466003402

```
In [39]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

[0.12455754]

8480.156871173602

```
In [40]: y_pred_elastic=regr.predict(X_train)
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print(mean_squared_error)
```

3708273.194830543